

High-Throughput, Low-Latency DA-FIR Filter with Hybrid Multiplier Architecture

*An Internship Report Submitted in the
Partial Fulfillment of the Requirements
for the Award of the completion of*

EPBL INTERNSHIP PROGRAM 2026

Submitted by

Anantha Akshay Kumar Reddy	23881A0406
Ayesha Siddiqua	23881A0407
Bhukya Anil	23881A0408

SUPERVISOR

Prof. S. Rajendar

Professor

VLSI Design and Systems

June, 2026

Abstract

This paper describes the construction of a high-throughput, low-latency Distributed Arithmetic Finite Impulse Response (DA-FIR) filter for noise removal in Software-Defined Radio (SDR), which incorporates an Approximate Karatsuba Multiplier (AKM) and a Variable Latency Carry Skip Adder (VLCSA). To address constraints in static truncation and rigid architectures, the design includes dynamic adaptive truncation, which adjusts the amount of truncated least significant bits (LSBs) in real time, as well as error compensation circuits to reduce truncation-induced distortions. A hybrid multiplier design combines AKM with Booth encoding to speed up partial product generation, whereas enhanced adder logic combines VLCSA with parallel prefix adders (e.g., Kogge-Stone) to accelerate crucial routes. FPGA-specific optimizations use block RAM (BRAM) and pipelining to improve resource consumption and clock frequency. Furthermore, machine learning-driven coefficient adaptation dynamically adjusts filter parameters to different noise profiles, while energy-efficient clock gating reduces idle power usage. The implementation focuses on real-world SDR situations, stressing adaptability, precision, and energy savings through modular hardware architecture.

Keywords: DA-FIR Filter, Dynamic Adaptive Truncation, Error Compensation Circuits, Block RAM.

Table of Contents

Title	Page No.
Abstract	i
List of Tables	v
List of Figures	vi
Abbreviations	vi
CHAPTER 1 Introduction	1
1.1 Introduction	1
1.2 Background and Motivation	2
1.3 Objectives of the Project Work	6
1.4 Organization of the Report	7
CHAPTER 2 Literature Survey	8
2.1 Introduction	8
2.2 Review of Prior Research	8
2.2.1 Distributed Arithmetic (DA) in FIR Filters	8
2.2.2 Approximate Karatsuba Multipliers	9
2.2.3 Variable-Length Carry Select Adders (VLCSA)	9
2.2.4 Modified FIR Filter and High-Speed Adder Design	10
2.2.5 FIR Filter Using Wallace Reduction Tree	11
2.3 Comparison: FIR vs. IIR Filters	12
2.3.1 FIR Filters	12
2.3.2 IIR Filters	12
2.3.3 Conclusion of Comparison	13
2.4 Types of FIR Filter Architectures	13
2.4.1 Direct Form FIR Filters	13
2.4.2 Transpose Form FIR Filters	13
2.4.3 Distributed Arithmetic (DA) FIR Filters	13
2.4.4 Parallel FIR Filters	14
2.5 Distributed Arithmetic-Based FIR Filters	14
2.5.1 Modified Distributed Arithmetic (MDA)	14
2.5.2 LUT-less DA	14
2.5.3 Partitioned LUT DA	14

2.6	Adder Architectures in FIR Filters	15
2.6.1	Ripple Carry Adder (RCA)	15
2.6.2	Carry Select Adder (CSLA)	15
2.6.3	Variable-Length Carry Select Adder (VLCSA)	15
2.7	Multiplier Architectures in FIR Filters	15
2.7.1	Booth and Baugh-Wooley Multipliers	15
2.7.2	Vedic Multiplier	15
2.7.3	Wallace Tree Multiplier	16
2.7.4	Approximate Karatsuba Multiplier (AKM)	16
2.8	Why AKM and VLCSA Are Preferred	16
2.8.1	AKM Advantages	16
2.8.2	VLCSA Advantages	16
2.9	Comparison Table	16
2.10	Identified Research Gaps	17
2.11	Performance Metrics Comparison	19
2.12	Summary	19
CHAPTER 3	Existing Work	21
3.1	Introduction	21
3.2	System Architecture Overview	21
3.3	Distributed Arithmetic: Mathematical Background	22
3.4	Approximate Karatsuba Multiplier: Logical Implementation	22
3.5	Variable-Length Carry Select Adder (VLCSA)	23
3.5.1	VLCSA Delay Model	23
3.6	Signal Flow and Control Strategy	23
3.7	Truth Table Derived from Simulation	24
3.8	Pipelining and Parallelism	24
3.9	Functional Block Summary	24
3.10	Control FSM Design	24
3.11	Advantages of Proposed Design	25
3.12	Conclusion	25
CHAPTER 4	Methodology	26
4.1	Introduction	26
4.2	System Architecture Overview	26
4.3	Distributed Arithmetic: Mathematical Background	28
4.4	Approximate Karatsuba Multiplier: Logical Implementation	29
4.5	Variable-Length Carry Select Adder (VLCSA)	29
4.5.1	VLCSA Delay Model	30
4.6	Signal Flow and Control Strategy	30
4.7	Truth Table Derived from Testbench	31

4.8	Pipelining and Parallelism	31
4.9	Functional Blocks Summary	32
4.10	Control FSM (Finite State Machine)	33
4.11	Advantages of Proposed Design	33
4.12	Conclusion	34
CHAPTER 5	Results and Discussion	35
5.1	Introduction	35
5.2	Overview of Results	35
5.3	Analysis and Interpretation	36
5.4	Evaluation of Quality Factors	36
5.5	Waveform Simulation Results	37
5.6	RTL Architecture Comparison	38
5.7	Conclusion	38
CHAPTER 6	Conclusions and Future Scope	42
6.1	Conclusions	42
6.2	Future Scope of Work	42
REFERENCES		44

List of Tables

2.1	Parallel FIR Filter Architecture Using Modified Adder Logic . .	11
2.2	Resource Utilization Comparison	11
2.3	Comparison of FIR Filter Architectures	17
2.4	Comparison of Adders and Multipliers in FIR Filters	19
3.1	Testbench-Derived Functional Table	24
4.1	Functional Table from Testbench Simulation	31

List of Figures

2.1	Approximate Karatsuba Multiplier Structure	9
2.2	VLCSA Structure Based on Brent-Kung Network	10
2.3	Parallel FIR Filter Architecture Using Modified Adder Logic . .	10
2.4	FIR Filter Architecture Using Wallace Reduction Tree and Vedic Multiplier	11
3.1	FIR Filter Architecture	22
4.1	DA-FIR Filter Design Using Approximate Karatsuba Multiplier and VLCSA	27
5.1	Simulation waveform of the earlier implementation	37
5.2	Simulation waveform of the proposed implementation	38
5.3	Approximate Karatsuba Multiplier – Earlier Design	39
5.4	VLCSA – Earlier Design	39
5.5	Complete Architecture – Earlier Design	39
5.6	Approximate Karatsuba Multiplier – Proposed Design	40
5.7	VLCSA – Proposed Design	40
5.8	Adder 4 Bit – Proposed Design	40
5.9	Complete Architecture – Proposed Design	40

Abbreviations

Abbreviation	Description
CMOS	Complementary Metal Oxide Semiconductor
FIR	Finite Impulse Response
IIR	Infinite Impulse Response
DA	Distributed Arithmetic
AKM	Approximate Karatsuba Multiplier
VLCSA	Variable-Length Carry Select Adder
LUT	Look-Up Table
SDR	Software-Defined Radio
MAC	Multiply-Accumulate
CSLA	Carry Select Adder
LFA	Ladner Fischer Adder
KSA	Kogge-Stone Adder
RCA	Ripple Carry Adder
HDL	Hardware Description Language
FPGA	Field Programmable Gate Array
ASIC	Application Specific Integrated Circuit
VLSI	Very Large Scale Integration
BRAM	Block Random Access Memory
IoT	Internet of Things
RTL	Register Transfer Level
FSM	Finite State Machine
KHz	Kilohertz
MHz	Megahertz

CHAPTER 1

Introduction

1.1 Introduction

Finite Impulse Response (FIR) filters play a critical role in digital signal processing applications such as audio enhancement, wireless communication, image processing, and biomedical signal analysis. Among various implementation techniques, Distributed Arithmetic (DA) offers an efficient hardware-centric approach by eliminating the need for traditional multipliers. However, conventional DA-FIR architectures often suffer from increased latency and limited throughput due to the bottlenecks in the accumulation and multiplication stages. The project titled "High-Throughput, Low-Latency DA-FIR Filter with Hybrid Multiplier Architecture" focuses on designing an efficient digital FIR filter using a Distributed Arithmetic (DA) approach enhanced with a hybrid multiplier setup. FIR filters are widely used in digital signal processing for applications such as audio filtering, image enhancement, and communication systems. However, traditional implementations often face challenges related to speed and hardware resource usage.

To address these issues, this project combines the DA technique with a Hybrid Multiplier Architecture that includes an Approximate Karatsuba Multiplier for faster multiplication and a Variable-Length Carry Select Adder (VLCSA) for efficient addition. The integration of these components aims to reduce latency and increase the overall throughput of the filter without significantly increasing area or power consumption.

This design is implemented in Verilog HDL and tested through simulation to verify its performance. The result is a high-speed, low-latency FIR filter architecture suitable for real-time and resource-constrained applications.

The primary aim of this project is to design and implement a high-performance FIR filter using Distributed Arithmetic (DA) enhanced by a

Hybrid Multiplier Architecture that delivers low latency and high throughput while maintaining hardware efficiency. This is achieved by integrating an Approximate Karatsuba Multiplier and a Variable-Length Carry Select Adder (VLCSA) into the filter structure to optimize computation time and reduce resource usage. The scope of the project includes:

- Designing a DA-based FIR filter suitable for real-time signal processing applications.
- Implementing a hybrid multiplier system combining approximate computation techniques (Karatsuba) with fast addition (VLCSA).
- Writing Verilog RTL code for simulation, synthesis, and validation of the design.
- Evaluating performance metrics such as delay, area, and power against traditional FIR architectures.
- Demonstrating the feasibility of using approximate arithmetic in DSP applications where a trade-off between speed and accuracy is acceptable.
- Laying the groundwork for future extensions, including FPGA/ASIC implementation and further optimization for low-power systems.

1.2 Background and Motivation

Modern embedded and edge-based DSP systems demand both speed and efficiency. Traditional FIR filter implementations using serial multiplication and accumulation limit performance, especially in high-order filters or systems requiring real-time computation. Distributed Arithmetic is known for its multiplier-less structure, using precomputed look-up tables and shift-accumulate logic. However, the LUT size grows exponentially with the number of taps, making it challenging for higher-order filters.

To enhance performance without significantly increasing area or power, hybrid multiplier strategies have been introduced. The Approximate Karatsuba Multiplier reduces partial product computations while maintaining acceptable

accuracy, and the VLCSA improves addition speed by handling carry propagation more efficiently than ripple-carry adders.

The motivation behind this project is to combine the benefits of DA architecture with a hybrid multiplication strategy to achieve a FIR filter design that offers low latency and high throughput without compromising significantly on accuracy—perfect for power-aware, performance-constrained applications. This project is situated in the field of Digital Signal Processing (DSP) and VLSI (Very-Large-Scale Integration) Design, specifically focusing on hardware-efficient implementations of FIR (Finite Impulse Response) filters. It combines concepts from digital electronics, computer architecture, and approximate computing to improve the performance of signal processing systems.

Within DSP, FIR filters are essential components used for shaping, modifying, and analyzing signals in various domains such as audio processing, image enhancement, telecommunications, biomedical instrumentation, and radar systems. Efficient hardware implementation of these filters is crucial for real-time processing and low-power embedded systems.

The project also involves high-level digital design using HDL (Hardware Description Language), particularly Verilog, and explores optimization techniques like Distributed Arithmetic (DA), approximate multiplication (using the Karatsuba algorithm), and advanced adder architectures (like VLCSA) to reduce delay, power consumption, and hardware complexity.

Traditionally, FIR filters were implemented using multiply-and-accumulate (MAC) units. However, as the filter order increased, the complexity and power consumption of these systems became a bottleneck, especially in portable and embedded devices.

To address this, researchers introduced Distributed Arithmetic (DA) in the 1970s, which replaced multipliers with precomputed lookup tables (LUTs) and shift-accumulate operations—greatly improving efficiency in hardware-based systems. While DA significantly reduced the use of multipliers, it introduced new challenges like LUT size growth and delay in accumulation, especially for high-order filters.

As technology advanced, focus shifted toward custom multiplier architectures. The Karatsuba algorithm, initially developed in the 1960s for fast

multiplication, was later adapted in approximate forms for hardware applications to reduce area and energy consumption. In parallel, carry-select and carry-lookahead adders were developed to speed up arithmetic operations, leading to innovations like the Variable-Length Carry Select Adder (VLCSA) for even better timing performance.

Combining these innovations, recent trends emphasize hybrid architectures—mixing approximate multipliers with fast adders and DA logic—to meet the performance demands of modern DSP applications. This project builds on these advancements to deliver a low-latency, high-throughput FIR filter, optimized for today’s real-time and power-sensitive systems.

This project builds upon several key research directions in Distributed Arithmetic, Approximate Computing, and Hybrid Multiplier Architectures:

Distributed Arithmetic (DA) FIR Filters: Distributed Arithmetic was first proposed as a memory-based technique to perform inner product operations without conventional multipliers. Research by White and Parks (1970s) demonstrated how DA can significantly reduce the complexity of FIR filters by using LUTs and shift-accumulate logic. Later improvements focused on optimizing LUT size and reducing delay through bit-serial and bit-parallel implementations.

Karatsuba-Based Multipliers: The Karatsuba algorithm, originally a divide-and-conquer method for multiplication, was introduced in the 1960s. Its hardware adaptation came much later, especially in approximate computing, where it helped reduce the number of partial products and improve speed. Recent studies explored Approximate Karatsuba Multipliers, which sacrifice minor accuracy for significant gains in area and power efficiency—making them ideal for DSP systems where exact precision isn’t always necessary.

Approximate Computing in DSP: Approximate computing has emerged as a valuable approach in error-tolerant applications such as multimedia and communication. Research has shown that introducing approximation in arithmetic units—like multipliers and adders—can drastically improve power consumption and speed while keeping output quality within acceptable ranges.

VLCSA (Variable-Length Carry Select Adder): Fast adder architectures, including Carry Select and Carry Lookahead Adders, have been

extensively studied for reducing propagation delay in arithmetic operations. VLCSA improves over fixed-length structures by dynamically adjusting adder block size based on operand significance, as explored in VLSI-oriented research. This makes it especially suitable for high-performance FIR filter designs.

Hybrid Architectures in FIR Design: Some recent works have proposed combining DA with custom multiplier and adder designs. For instance, hybrid architectures incorporating approximate multipliers and high-speed adders have shown promising results in terms of speed, area, and power trade-offs. These designs provide a solid foundation and motivation for your project, which aims to push these improvements further.

As digital systems become faster and more integrated into everyday devices—from smartphones to medical wearables—there is an increasing demand for real-time digital signal processing (DSP) solutions that are both high-performance and energy-efficient. Among the most commonly used DSP components are FIR filters, which are crucial for signal conditioning, noise reduction, and feature extraction.

While Distributed Arithmetic (DA) offers an efficient hardware approach by minimizing the use of multipliers, traditional DA implementations still suffer from limited throughput and high latency, especially as filter length increases. Additionally, the growing complexity of applications in IoT, AI edge devices, and 5G systems calls for optimized hardware that can maintain speed without significantly increasing area or power.

This project addresses this critical gap by:

- DA with a hybrid multiplier-adder architecture, merging Approximate Karatsuba Multipliers with VLCSA to accelerate processing.
- Reducing latency and improving throughput while maintaining acceptable accuracy levels—ideal for time-critical and energy-constrained applications.
- Providing a scalable and resource-efficient solution that can be adapted for high-order filters or integrated into custom VLSI systems.

At the same time, recent advancements in approximate computing and custom

multiplier designs (like Approximate Karatsuba Multipliers) have shown potential to meet these performance demands—but haven’t yet been fully integrated into a DA-based FIR architecture in a standardized or streamlined way.

1.3 Objectives of the Project Work

- To design and implement a Distributed Arithmetic (DA)-based FIR filter that eliminates the need for conventional multipliers, optimizing for real-time signal processing applications.
- To integrate a Hybrid Multiplier Architecture by combining an Approximate Karatsuba Multiplier and a Variable-Length Carry Select Adder (VLCSA) to enhance computation speed and reduce latency.
- To achieve high throughput by minimizing the critical path delay in the filter architecture using optimized arithmetic operations.
- To reduce overall power consumption and hardware complexity, making the design suitable for low-power and resource-constrained environments like IoT and embedded systems.
- To develop Verilog RTL code for the proposed design and simulate it to validate the filter’s functionality, accuracy, and performance.
- To analyze key performance metrics such as delay, area, power, and throughput, and compare them against conventional FIR filter architectures.
- To explore the impact of approximation techniques on the filter’s output accuracy and assess their suitability in practical DSP applications.

The primary objective of this project is to design and implement a high-performance Finite Impulse Response (FIR) filter utilizing Distributed Arithmetic (DA) combined with a hybrid multiplier architecture. This architecture integrates an Approximate Karatsuba Multiplier and a Variable-Length Carry Select Adder (VLCSA) to enhance computational efficiency. The project aims

to achieve low latency and high throughput while minimizing hardware complexity and power consumption. Implementation will be carried out using Verilog HDL, with thorough simulation and analysis to validate functionality and performance metrics such as speed, area, and power efficiency. Additionally, the project will assess the impact of approximation techniques on filter accuracy to ensure suitability for real-time digital signal processing applications.

1.4 Organization of the Report

The final section of **Chapter 1** provides an outline of the structure of the entire report. It explains the organization and flow of the chapters and gives the reader an understanding of how the content will be presented. This section can be a simple list or paragraph summarizing each chapter's main points.

Chapter 2: Literature Review A review of existing research and developments related to FIR filters, Distributed Arithmetic, approximate multipliers, and adder architectures is presented to establish the foundation and context for the current project.

Chapter 3: Design of Distributed Arithmetic FIR Filter This chapter details the design approach of the DA-based FIR filter, discussing the theoretical underpinnings and practical considerations in implementing the filter architecture.

Chapter 4: Implementation of Hybrid Multiplier Architecture This section elaborates on the integration of the Approximate Karatsuba Multiplier and Variable-Length Carry Select Adder (VLCSA) within the DA-FIR filter framework, highlighting the rationale behind the chosen design strategies.

Chapter 5: Results and Discussion Simulation results are presented and analyzed to evaluate the performance of the proposed filter design. Metrics such as speed, area, power consumption, and accuracy are examined and compared with traditional FIR filter implementations.

Chapter 6: Conclusions and Recommendations This chapter summarizes the findings, highlighting the achievements and limitations of the project.

CHAPTER 2

Literature Survey

2.1 Introduction

Finite Impulse Response (FIR) filters are fundamental components in digital signal processing (DSP), widely used in applications such as audio processing, communications, and image filtering. Traditional implementations of FIR filters rely heavily on Multiply-Accumulate (MAC) operations, which can be resource-intensive and power-hungry. To address these challenges, alternative architectures have been explored, notably Distributed Arithmetic (DA), Approximate Karatsuba Multipliers, and Variable-Length Carry Select Adders (VLCSA). These approaches aim to optimize performance, reduce power consumption, and maintain accuracy, making them suitable for high-throughput, low-latency applications [1].

2.2 Review of Prior Research

2.2.1 Distributed Arithmetic (DA) in FIR Filters

DA is a technique that replaces traditional multipliers in FIR filters with precomputed lookup tables (LUTs), shift registers, and adders, leading to reduced hardware complexity and power consumption. Studies have demonstrated that DA-based FIR filters offer significant improvements in area, power, and throughput compared to conventional implementations. For instance, research by Grande Naga Jyothi and Sridevi Sriadibhatla presents various DA architectures, including LUT-based and LUT-less designs, highlighting their advantages in terms of area and power efficiency.

2.2.2 Approximate Karatsuba Multipliers

The Karatsuba algorithm, a divide-and-conquer method for multiplication, has been adapted into approximate forms to enhance performance in error-resilient applications. Approximate Karatsuba Multipliers reduce hardware requirements and latency by simplifying multiplication operations, making them suitable for applications where exact precision is not critical. Comparative analyses have demonstrated that these multipliers offer improvements in accuracy, latency, and hardware efficiency over traditional multipliers [2].




Figure 2.1: Approximate Karatsuba Multiplier Structure

2.2.3 Variable-Length Carry Select Adders (VLCSA)

Carry Select Adders (CSAs) are known for their speed in arithmetic operations. The Variable-Length Carry Select Adder (VLCSA) enhances this by dynamically adjusting the block sizes based on operand significance, leading to further reductions in propagation delay. Studies have shown that VLCSAs offer a balance between speed and area efficiency, making them advantageous in high-performance computing scenarios.




Figure 2.2: VLCSA Structure Based on Brent-Kung Network

2.2.4 Modified FIR Filter and High-Speed Adder Design

This paper focuses on replacing standard adders with high-speed variants like Modified Carry Select Adders (CSLA), Ladner Fischer Adders (LFA), and Brent-Kung adders in a parallel FIR filter.

Highlights include:

- Parallel prefix computation reduces delay significantly.
- D-latch based CSLA shows a 23% decrease in propagation delay.
- Power reduction from 139 mW (standard CSLA) to 110 mW.




Figure 2.3: Parallel FIR Filter Architecture Using Modified Adder Logic

Table 2.1: Parallel FIR Filter Architecture Using Modified Adder Logic


Adder Type	Delay (ns)	Power (mW)
Standard CSLA	25.18	139
D-Latch Based CSLA	16.71	110
LFA	12.40	93

2.2.5 FIR Filter Using Wallace Reduction Tree

This paper introduces a high-speed FIR filter using a combination of Wallace Tree Reduction and Vedic Multipliers for fast multiplication and accumulation.[**biblatex-control**]

Key aspects:

- Vedic Multipliers enable parallel generation of partial products.
- Wallace Tree reduces critical path delay by adding in a tree-like structure.
- Achieved 57% LUT reduction and 50% register savings over traditional FIR filters.

**Figure 2.4:** FIR Filter Architecture Using Wallace Reduction Tree and Vedic Multiplier**Table 2.2:** Resource Utilization Comparison

Architecture	LUT Usage	Slice Registers
Baugh-Wooley FIR	5000	3200
Wallace-Vedic FIR	2150	1580

2.3 Comparison: FIR vs. IIR Filters

2.3.1 FIR Filters

FIR filters are characterized by a finite impulse response, meaning that their output settles to zero in a finite amount of time after the input signal ends. This property ensures inherent stability, as the filters do not rely on feedback components. FIR filters are non-recursive and utilize only current and past input samples in their calculations. This architecture makes them ideal for systems that require linear phase response, such as audio and biomedical signal processing, where phase distortion must be avoided.[bibtex-control]

- **Stability:** FIR filters are always stable due to the absence of feedback loops.
- **Linear Phase:** They can be designed to have linear phase, which means all frequency components are delayed equally, preserving waveform shape.
- **Implementation Simplicity:** Their structure is straightforward, making them easier to design and test.
- **Resource Intensive:** Achieving sharp frequency cutoffs requires more filter coefficients, increasing computational and memory demands.

2.3.2 IIR Filters

IIR filters utilize both past inputs and past outputs (feedback) to compute the current output. This recursive behavior enables them to achieve a given frequency response with fewer coefficients compared to FIR filters, thereby reducing computational requirements. However, the presence of feedback also introduces the possibility of instability if not properly designed.

- **Efficiency:** IIR filters typically require fewer coefficients for the same filtering effect.
- **Non-linear Phase:** Most IIR filters do not offer linear phase response, which can distort the signal waveform.

- **Stability Concerns:** Improperly designed IIR filters can become unstable.
- **Design Complexity:** The recursive structure adds complexity in terms of analysis and implementation.

2.3.3 Conclusion of Comparison

FIR filters are widely preferred in applications requiring predictable phase response and guaranteed stability. Despite higher computational costs, their deterministic behavior and ease of implementation in hardware (especially with techniques like DA) make them highly suitable for FPGA and ASIC designs.

2.4 Types of FIR Filter Architectures

2.4.1 Direct Form FIR Filters

This is the most straightforward implementation involving a cascade of Multiply-Accumulate (MAC) operations. Each input sample is multiplied by a coefficient, and the results are summed to produce the output. While this form is easy to understand and implement, it is not optimized for hardware due to the sequential nature of the MAC operations.

2.4.2 Transpose Form FIR Filters

Transpose form rearranges the direct form architecture to be more amenable to pipelining. This form allows for easier implementation of high-speed filters in hardware by reducing the critical path through increased parallelism.

2.4.3 Distributed Arithmetic (DA) FIR Filters

DA replaces multipliers with precomputed values stored in lookup tables (LUTs), shift registers, and adders. It is especially useful when filter coefficients are fixed, allowing for extensive reuse of LUT values. This approach reduces the number of arithmetic operations significantly.

2.4.4 Parallel FIR Filters

In this form, the FIR filter is partitioned to process multiple inputs or outputs in parallel. This architecture enhances throughput but requires additional hardware resources. It is often employed in systems with high sample rate requirements.

2.5 Distributed Arithmetic-Based FIR Filters

2.5.1 Modified Distributed Arithmetic (MDA)

MDA introduces optimization by reusing parts of the hardware. Multiplexer-based designs allow shared logic blocks, effectively reducing the memory size and circuit area. In practice, MDA implementations have demonstrated up to 42% reduction in cell count and about 2% reduction in power consumption, making it suitable for power-sensitive applications such as wearable biomedical devices.

2.5.2 LUT-less DA

To overcome the exponential growth of LUTs with filter order, LUT-less architectures use algorithmic computation during runtime. Although this may introduce a slight delay penalty, the substantial area and power savings often justify the trade-off. These architectures are more scalable and are better suited for dynamically reconfigurable systems.

2.5.3 Partitioned LUT DA

This method divides large LUTs into smaller, manageable segments, allowing for more efficient memory usage. It significantly reduces access time and is particularly useful in high-tap filters, where large coefficient sets are involved. Partitioning also enables selective updating of LUT contents, providing more flexibility.

2.6 Adder Architectures in FIR Filters

2.6.1 Ripple Carry Adder (RCA)

RCA is the simplest adder design, where each bit waits for the carry from the previous bit to propagate. While easy to implement, it suffers from significant delay, especially as the operand size increases.

2.6.2 Carry Select Adder (CSLA)

CSLA improves speed by computing two sums in parallel (one assuming carry-in is 0 and another assuming it is 1), and then selecting the correct result using a multiplexer. Although faster than RCA, it incurs higher area overhead due to duplication of logic.

2.6.3 Variable-Length Carry Select Adder (VLCSA)

VLCSA is an advanced form of CSLA that uses variable-sized blocks to balance the area-delay trade-off. It minimizes unnecessary duplication and uses techniques such as Brent-Kung networks to optimize carry propagation. This makes VLCSA highly suitable for real-time DSP applications, where both speed and resource efficiency are critical.

2.7 Multiplier Architectures in FIR Filters

2.7.1 Booth and Baugh-Wooley Multipliers

These are traditional methods for signed multiplication. Booth encoding reduces the number of partial products, and Baugh-Wooley allows efficient signed operations. However, both require significant hardware and are relatively slower compared to modern alternatives.

2.7.2 Vedic Multiplier

Inspired by ancient Indian mathematics, Vedic multipliers offer high-speed multiplication by using parallelism in partial product generation. They are ef-

fective for small operand sizes and are often used in Wallace tree configurations for high-speed DSP blocks.

2.7.3 Wallace Tree Multiplier

Wallace tree structures allow parallel reduction of partial products using layers of compressors. This results in low latency and is ideal for large bit-width multiplications.[**biblatex-control**]

2.7.4 Approximate Karatsuba Multiplier (AKM)

AKM employs a recursive divide-and-conquer approach and introduces approximation to simplify intermediate calculations. This reduces the number of logic gates, power consumption, and overall delay. While there is a trade-off in accuracy, AKM is highly efficient for error-tolerant applications like image filtering and certain machine learning models.

2.8 Why AKM and VLCSA Are Preferred

2.8.1 AKM Advantages

AKM is particularly well-suited for FIR filters due to its modularity and speed. It effectively reduces multiplication complexity by breaking down operations and reusing computations. Its approximation feature is controllable, allowing designers to fine-tune performance versus accuracy.

2.8.2 VLCSA Advantages

VLCSA addresses the critical path delay associated with large-bit addition. By partitioning adders into variable-length blocks, it provides a better balance of speed and area. Furthermore, it can be adapted to specific operand bit-widths, making it flexible for a wide range of DSP applications.

2.9 Comparison Table

Table 2.3: Comparison of FIR Filter Architectures

Feature	DA + AKM + VLCSA	Modified FIR with Adders	FIR with Wallace Tree
Architecture Type	Distributed Arithmetic	Parallel FIR Design	Convolution with Wallace Tree
Multiplier Technique	Approx. Karat-suba	Compressor-based	Vedic + Wallace
Adder Type	VLCSA (Brent-Kung)	CSLA, LFA, KSA	Ripple Carry in Wallace
Power Consumption	Low	Moderate to Low	Low
Delay Optimization	Excellent	Good	Excellent
LUT Efficiency	High	Moderate	Very High (57% reduced)
FPGA Compatibility	High (Xilinx)	High	High (No DSP48 used)

2.10 Identified Research Gaps

Despite notable advances in FIR filter design using Distributed Arithmetic, modified adder logic, and Wallace Tree multipliers, several open challenges remain that hinder achieving optimal performance for real-time, power-sensitive applications:

- **Limited Integration of Techniques:** Most research focuses on improving individual components (e.g., adders, multipliers, DA architecture) in isolation. There is a lack of holistic frameworks that integrate DA with approximate computing techniques (like AKM) and high-speed adders (like VLCSA or LFA). Such integration could synergistically reduce delay, area, and power, but is currently underexplored.[3]
- **Scalability to Higher-Order Filters:** As filter tap lengths increase, traditional DA-based designs suffer from exponential LUT growth. While truncation methods (e.g., LSB dropping) help reduce LUT size, they

introduce non-linear errors. There is a research need to develop scalable architectures that maintain accuracy while managing memory complexity in high-order designs.[4]

- **Error Management in Approximate Computing:** While approximate multipliers (e.g., Karatsuba) reduce latency and power, they also introduce noise and quantization errors. Most existing works lack rigorous analysis of how these errors propagate through FIR filters. Future designs should incorporate formal error-bounded models or adaptive correction mechanisms to balance accuracy and efficiency.[5]
- **Underutilization of Parallelism in FPGA Resources:** Though Wallace Trees and parallel FIR topologies exploit some degree of parallelism, many designs still do not fully leverage modern FPGA features such as dual-port RAM, DSP slices, or reconfigurable logic blocks. Efficient mapping and scheduling of operations to FPGA resources remain an open area for optimization.
- **Latency vs. Throughput Trade-offs:** Pipelining techniques and adder tree optimizations (e.g., Brent-Kung) are often used to reduce critical path delay, but may increase overall latency or resource usage. There is a need for customizable architectures that allow designers to tune performance metrics dynamically, based on the application requirements.
- **Lack of Benchmarking Against Industrial Standards:** Many experimental FIR architectures are evaluated only on basic Verilog simulations or on isolated Xilinx toolchains. There is a lack of standard benchmarking across real-world signal processing workloads or standardized datasets. Future work should include real-time audio, biomedical, or wireless communication benchmarks for comprehensive evaluation.

These gaps highlight the need for a new generation of FIR filter architectures that combine the strengths of existing methods with intelligent approximation, parallelism, and resource-aware optimization.

2.11 Performance Metrics Comparison

Table 2.4: Comparison of Adders and Multipliers in FIR Filters

Architecture	Delay (ns)	Power (mW)
Ripple Carry Adder	High (Slow)	Low
CSLA	Moderate	High
VLCSA	Low (Fast)	Moderate
Booth Multiplier	Moderate	High
Wallace Tree + Vedic	Low	Moderate
AKM	Low	Low

2.12 Summary

The review of literature reveals the breadth of innovation in FIR filter design, emphasizing three core directions: Distributed Arithmetic-based designs, adder-optimized parallel filter structures, and Wallace Tree-based multiplication with Vedic logic. Each approach has demonstrated measurable benefits in speed, area, or power efficiency — but also presents inherent trade-offs.

Distributed Arithmetic significantly reduces hardware multipliers and is well-suited for low-power DSP systems, especially when LUT growth is managed via truncation techniques. The integration of Approximate Karatsuba Multipliers further compresses computation complexity, making them ideal for energy-constrained or error-tolerant environments such as audio or image processing.

Meanwhile, modified adder architectures like CSLA, Ladner Fischer, and Brent-Kung offer fast carry propagation and minimal area overhead. Their effectiveness in parallel FIR filters proves useful in real-time systems, where high throughput and clock speed are essential.

The Wallace Tree-based architecture, enhanced by Vedic multipliers, introduces a powerful combination for high-speed filtering with efficient hardware utilization. These designs excel in reducing LUTs and slice registers, making them suitable for compact, high-performance FPGA implementations in IoT and biomedical signal analysis.

Collectively, these studies provide a strong technical foundation but also expose unaddressed intersections — namely, the lack of hybrid, scalable, and accuracy-aware architectures. The proposed research aims to bridge this gap by developing a composite FIR filter architecture that leverages DA, AKM, VLCSA, and Wallace Tree principles, thereby balancing speed, accuracy, power, and FPGA resource efficiency. This work aspires to contribute a more holistic, modular, and application-adaptive FIR design paradigm to the field of digital signal processing.

CHAPTER 3

Existing Work

3.1 Introduction

This chapter outlines the design methodology for a high-efficiency Finite Impulse Response (FIR) filter. The architecture integrates Distributed Arithmetic (DA), Approximate Karatsuba Multiplication (AKM), and Variable-Length Carry Select Addition (VLCSA) to deliver a performance-optimized solution suitable for real-time signal processing systems. The design is implemented in Verilog HDL and tested through a custom testbench to ensure correctness and performance.

The motivation behind this approach stems from the need to optimize arithmetic-heavy modules for embedded platforms, where power and area efficiency are critical. Traditional MAC-based FIR filters often suffer from high hardware overhead due to multipliers and long carry propagation delays. This architecture addresses these challenges using DA to eliminate runtime multipliers, AKM for fast and compact multiplication, and VLCSA to accelerate addition stages [6].

3.2 System Architecture Overview

The FIR filter architecture includes the following pipelined stages:

- Input Register Bank
- Distributed Arithmetic (DA) LUT Generator
- Approximate Karatsuba Multiplier (AKM)
- Variable-Length Carry Select Adder (VLCSA)
- Output Register




Figure 3.1: FIR Filter Architecture

The architecture is synchronized using a global clock signal, enabling each stage to process data concurrently. Once the pipeline is filled, new input samples are processed on each clock cycle, maximizing throughput.

3.3 Distributed Arithmetic: Mathematical Background

Distributed Arithmetic simplifies the dot product calculation by replacing real-time multiplications with look-up tables (LUTs) and shift-add operations. [1] Given a standard FIR filter equation:

3.4 Approximate Karatsuba Multiplier: Logical Implementation

AKM is based on the divide-and-conquer strategy of the Karatsuba algorithm. It reduces multiplication complexity by splitting inputs into high and low halves:

The product is calculated as:

$$P = X_H Y_H \cdot 2^n + ((X_H + X_L)(Y_H + Y_L) - X_H Y_H - X_L Y_L) \cdot 2^{n/2} + X_L Y_L \quad (3.1)$$

In the approximate version, cross terms such as X_{HY_L} and X_{LY_H} may be simplified or omitted to reduce logic, trading a small amount of precision for significant area and power savings.

3.5 Variable-Length Carry Select Adder (VLCSA)

VLCSA improves addition speed by precomputing results for carry-in values 0 and 1 across variable-sized blocks. A multiplexer then selects the correct output, reducing delay compared to ripple carry or fixed-block CSLA implementations [7].

3.5.1 VLCSA Delay Model

$$T_{VLCSA} = \max(t_{rca1}, t_{rca2}, \dots, t_{rcan}) + t_{mux} \cdot (n - 1) \quad (3.2)$$

3.6 Signal Flow and Control Strategy

Data flows through the filter as follows:

- Input samples are loaded into shift registers.
- Binary decomposed inputs index DA LUTs.
- Partial products are approximated via AKM.
- VLCSA accumulates these products with minimal delay.
- Output is captured in a register on the next clock cycle.

Key control signals:

- **clk** – Global clock
- **rst** – Reset
- **en** – Enable signal
- **in_data** – Serial input stream
- **out_data** – Final output result

Table 3.1: Testbench-Derived Functional Table

Cycle	Input Sample	DA LUT Output	AKM Product	Final Output
1	0001	$\$h_0\$$	$\$P_0\$$	$\$h_0\$$
2	0010	$\$h_1\$$	$\$P_1\$$	$\$h_0 + h_1\$$
3	0100	$\$h_2\$$	$\$P_2\$$	$\$h_0 + h_1 + h_2\$$
4	1000	$\$h_3\$$	$\$P_3\$$	$\$h_0 + h_1 + h_2 + h_3\$$

3.7 Truth Table Derived from Simulation

3.8 Pipelining and Parallelism

The filter is deeply pipelined across the DA, AKM, and VLCSA stages. Each stage works in parallel, improving throughput by overlapping operations. This ensures a new result is output each clock cycle after the initial latency [8].

Parallelism Details:

- DA operates on all input bits concurrently.
- AKM computes partial products in parallel.
- VLCSA precomputes for both carry-in values.

3.9 Functional Block Summary

- **Input Shift Register:** Stores prior samples.
- **LUT Generator:** Outputs partial sums.
- **AKM Unit:** Approximates multiplication.
- **VLCSA Tree:** Fast accumulation.
- **Final Register:** Captures output result.

3.10 Control FSM Design

The design uses a four-state FSM:

1. **IDLE:** Await enable.

2. **LOAD:** Load sample.
3. **COMPUTE:** Process through DA, AKM, and VLCSA.
4. **WRITE:** Store final output.

FSM transitions are edge-triggered and governed by clock and control signals to maintain pipeline synchronization.

3.11 Advantages of Proposed Design

- **High Throughput:** Pipeline allows one output per cycle.
- **Low Power:** Approximate multiplication and fewer transitions.
- **Reduced Delay:** VLCSA minimizes critical path.
- **Scalable:** Modular design supports varying filter sizes.

3.12 Conclusion

This methodology demonstrates the practical integration of DA, AKM, and VLCSA into a single FIR filter architecture. The design balances speed, power, and area efficiency, verified through simulation and waveform inspection. Its modularity and performance make it ideal for embedded DSP applications.

CHAPTER 4

Methodology

4.1 Introduction

This chapter presents the design methodology of a high-performance, low-power Finite Impulse Response (FIR) filter that integrates Distributed Arithmetic (DA), Approximate Karatsuba Multiplication (AKM), and Variable-Length Carry Select Addition (VLCSA). These optimizations are aimed at enhancing performance in real-time digital signal processing systems. The design is implemented in Verilog Hardware Description Language (HDL) and validated using a thorough verification environment [8].

The core motivation behind this architecture is the growing need for signal processing blocks that can function efficiently within resource-constrained embedded systems while delivering high-speed operations. Traditional Multiply-Accumulate (MAC) based filters, although effective, suffer from limitations related to multiplier hardware overhead and propagation delay in wide adders. This challenge is addressed through DA, which eliminates the need for run-time multipliers [2].

Additionally, AKM and VLCSA further contribute to computational savings. AKM offers approximate, fast multiplications using divide-and-conquer principles, while VLCSA significantly reduces adder delay by adjusting carry select stages dynamically. Together, these three techniques offer a synergistic enhancement to FIR filter architectures [bibtex-control].

4.2 System Architecture Overview

The system architecture is structured to incorporate pipeline stages for enhanced throughput while minimizing the logic required for arithmetic operations. It begins with an input register bank, which stores the incoming digital samples. These samples are processed using DA-based LUTs, producing partial

results depending on the binary decomposition of the input. The AKM blocks then approximate the multiplications of coefficients and inputs to reduce logic usage [3].

Following this, the partial products are passed to a VLCSA adder tree, which accumulates results with minimal propagation delay. The output from the adder tree is fed into an output register, which holds the filtered result. The design is synchronous and clock-driven, allowing for easy integration into larger pipelined systems.




Figure 4.1: DA-FIR Filter Design Using Approximate Karatsuba Multiplier and VLCSA

Each component is carefully modularized, facilitating easy synthesis, debugging, and future expansion. The pipelined nature of the architecture ensures that one input sample is processed every clock cycle after the pipeline is filled, making this design particularly effective for high-data-rate applications such as audio, image, and wireless signal processing.

The architectural blocks are:

- **Input Register Bank:** Buffers the input samples required for FIR operations.
- **DA-LUT Generator:** Computes bit-wise partial sums based on the binary input bits.
- **Karatsuba Partial Multiplier:** Efficiently handles large bit-width co-efficient multiplication.
- **VLCSA Accumulator:** Adds partial products with variable block widths for reduced delay.
- **Output Register:** Stores the final computed result.

4.3 Distributed Arithmetic: Mathematical Background

Distributed Arithmetic provides a technique for computing the dot product without real-time multipliers. The inputs are represented in bit-sliced binary form and processed using shift-add operations combined with LUTs containing precomputed weighted sums. This reduces the complexity of FIR filters by simplifying the multiplication task to a series of table lookups and additions. Given:

$$y[n] = \sum_{i=0}^{N-1} h[i] \cdot x[n-i] \quad (4.1)$$

This is rewritten using bit-sliced inputs $x[n]$ represented in binary:

$$y[n] = \sum_{k=0}^{B-1} 2^{-k} \cdot \left(\sum_{i=0}^{N-1} h[i] \cdot b_k[n-i] \right) \quad (4.2)$$

Where $b_k[n]$ is the k -th bit of the input word. The inner summation is precomputed and stored in a lookup table indexed by bit patterns.

This approach significantly reduces area in FPGA and ASIC designs. The LUTs are small for short filter lengths and can be partitioned or compressed

further using optimization techniques. DA is especially effective when coefficients are constant or change infrequently.

4.4 Approximate Karatsuba Multiplier: Logical Implementation

The Karatsuba algorithm is a divide-and-conquer strategy for fast multiplication. It splits the input into high and low halves and reduces the total number of multiplications required. Let X and Y be n -bit numbers:

$$\begin{aligned} X &= X_H \cdot 2^{n/2} + X_L \\ Y &= Y_H \cdot 2^{n/2} + Y_L \end{aligned}$$

Then:

$$P = X_H Y_H \cdot 2^n + ((X_H + X_L)(Y_H + Y_L) - X_H Y_H - X_L Y_L) \cdot 2^{n/2} + X_L Y_L \quad (4.3)$$

In the approximate version, certain intermediate products like $X_H Y_L$ and $X_L Y_H$ are truncated or simplified to save logic.

4.5 Variable-Length Carry Select Adder (VLCSA)

VLCSA is a speed-optimized adder that improves upon the basic Carry Select Adder by allowing variable-length block partitioning. Instead of having fixed-sized adder blocks, VLCSA adjusts the block sizes based on operand significance, ensuring that more critical paths receive faster computation.

This block-wise structure allows simultaneous computation of sum outputs for possible carry-in values (0 and 1). The final result is selected using a multiplexer, which introduces a small delay but removes the sequential carry chain found in ripple adders. The flexibility in block size improves delay optimization by adapting to the specific operand size and propagation characteristics [9].

In this FIR filter design, VLCSA is used in the accumulation stage where

timing is critical. Its ability to process sums faster than RCA or standard CSLA results in improved filter throughput without drastically increasing hardware cost. This makes VLCSA a key contributor to overall performance enhancement.

4.5.1 VLCSA Delay Model

Let t_{mux} be the delay of a multiplexer, and t_{rca} be the delay of a ripple adder in each block:

$$T_{VLCSA} = \max(t_{rca1}, t_{rca2}, \dots, t_{rcan}) + t_{mux} \cdot (n - 1) \quad (4.4)$$

4.6 Signal Flow and Control Strategy

Signal flow begins with loading the input samples into the shift register bank. On each clock pulse, new samples are shifted in while older ones are retained. This data is then fed into the DA logic, which uses the binary representation of inputs to address LUT entries. The selected values are weighted and shifted according to bit significance[4].

The next step involves multiplication of precomputed outputs with filter coefficients using AKM. These approximated multiplications reduce the critical path and power dissipation. The products are fed to the VLCSA unit for accumulation. Because of pipelined architecture, accumulation results are produced every clock cycle after initialization.

Control signals such as `clk`, `rst`, and `en` manage the flow of data. A reset initializes all registers, while enable allows loading of new data. This synchronous design ensures reliable and deterministic behavior, crucial for embedded and real-time applications.

The design operates on clock edge synchronization. Inputs are loaded on the positive edge, and partial sums are updated each cycle through pipelined registers. Enable signals control register loading and LUT access. Reset initializes the system [5].

- **clk:** Master clock signal.

- **rst:** System reset (active high).
- **en:** Control signal to load data.
- **in_data:** Serial input sample stream.
- **out_data:** Filtered output result.

The accuracy of the output was further validated through waveform inspection using a simulator. For each input sequence, the resulting output matched expected FIR values. Corner cases such as all-zeros and all-ones were also tested and confirmed stable operation.

These results confirm that the design maintains accuracy across different data inputs and successfully integrates all three computational components. The testbench also includes functional assertions and delays to simulate real-world scenarios.

4.7 Truth Table Derived from Testbench

Using the testbench simulation results, we derive a set of test cases for the filter.

Table 4.1: Functional Table from Testbench Simulation

Cycle	Input Sample	DA LUT Output	AKM Product	Final Output
1	0001	h_0	P_0	h_0
2	0010	h_1	P_1	$h_0 + h_1$
3	0100	h_2	P_2	$h_0 + h_1 + h_2$
4	1000	h_3	P_3	$h_0 + h_1 + h_2 + h_3$

4.8 Pipelining and Parallelism

To enhance performance, the architecture uses pipelining in key stages—DA LUT computation, AKM multiplication, and VLCSA accumulation. Each stage processes data in parallel, enabling the next input sample to be accepted even before the previous result is fully output.

This overlapping of operations significantly increases throughput. For a filter with depth d , results are produced every clock cycle after the initial d

cycles. The pipelining mechanism is supported by flip-flop based intermediate registers and clock enable signals that synchronize data propagation.

Parallelism is further achieved in VLCSA, where computations for both carry-in conditions are performed simultaneously. This reduces critical path delay and ensures fast accumulation. DA logic also operates in parallel bit planes, improving LUT access speed.

Parallel computation is used in:

- VLCSA blocks computing for carry-in 0 and 1.
- DA table output based on all active input bits.
- Karatsuba partial products.

4.9 Functional Blocks Summary

Each functional unit is designed to execute specific roles in the filter pipeline. The input register bank ensures consistent availability of past samples, critical for FIR filter operation. The DA LUT generator translates binary input to precomputed sums using efficient memory addressing.

The AKM multiplier executes fast, low-precision multiplications that help reduce power usage. Each product is then added using VLCSA to generate the total sum. This unit is specifically optimized for minimal carry propagation [10].

Finally, the output register captures the complete result and holds it for access or further processing. All these blocks are interconnected using control signals and data buses, providing a clean and modular design suited for FPGA or ASIC implementation.

- **Input Shift Register:** Stores the last N samples.
- **LUT Generator:** Outputs precomputed values based on input bit combination.
- **AKM Unit:** Performs approximated coefficient multiplications.
- **VLCSA Tree:** Adds the results quickly.

- **Final Register:** Stores result per clock cycle.

4.10 Control FSM (Finite State Machine)

The FSM manages the sequencing of operations across various modules. It consists of four primary states—IDLE, LOAD, COMPUTE, and WRITE. Each state is governed by system control signals and moves to the next upon clock edge or status flag changes [10].

In IDLE, the system awaits enablement. Once active, LOAD fetches the current sample into registers. COMPUTE triggers DA processing, AKM multiplication, and VLCSA accumulation. Upon completion, WRITE stores the result to the output register.

This FSM-based control ensures deterministic operation and makes the design easier to test and verify. FSM transitions also help synchronize internal pipeline stages with external data interface timing. The FSM transitions through the following states:

1. **IDLE:** Awaiting enable signal.
2. **LOAD:** Load current input sample.
3. **COMPUTE:** Generate DA, AKM, and VLCSA results.
4. **WRITE:** Store final output.

4.11 Advantages of Proposed Design

The proposed architecture offers multiple advantages over conventional FIR filters. By replacing run-time multipliers with DA LUTs and AKM, hardware complexity and power usage are greatly reduced. The use of VLCSA ensures fast accumulation, which is often a bottleneck in arithmetic pipelines.

Another major benefit is modularity. Each block is designed to be self-contained, making the overall system scalable. This modularity also eases verification and hardware debugging processes. Pipelining ensures that once the pipeline is filled, new results are produced every clock cycle [11].

Finally, the architecture is highly configurable. The number of filter taps, LUT size, and approximation accuracy can be tuned based on application needs. This adaptability makes it suitable for a wide range of DSP tasks including audio filtering, radar processing, and wearable biomedical systems.

- **High Throughput:** Achieved through pipelining and parallel blocks.
- **Low Power:** Approximate multipliers reduce switching activity.
- **Scalability:** Modular design supports expansion to larger tap filters.
- **Reduced Delay:** VLCSA minimizes critical path.

4.12 Conclusion

This chapter has outlined the detailed methodology behind the design of a high-speed, low-power FIR filter using Distributed Arithmetic, Approximate Karatsuba Multipliers, and Variable-Length Carry Select Adders. Each subsystem is crafted to minimize delay and maximize efficiency, resulting in an architecture ideal for real-time signal processing environments.

Through mathematical modeling, truth tables, architectural breakdown, and simulation-backed validation, we demonstrate that the proposed methodology delivers on performance and resource efficiency. With its modular design and scalable architecture, this filter framework serves as a solid foundation for future enhancements and practical deployments in embedded DSP systems.

CHAPTER 5

Results and Discussion

5.1 Introduction

The evaluation of digital systems such as FIR filters must be thorough to ensure the design meets critical specifications such as timing, area, power, and functional correctness. This chapter presents the measured outcomes of the proposed FIR filter using synthesis, waveform simulation, and post-implementation design reports.

The analysis covers both an earlier design and the optimized version that incorporates Distributed Arithmetic (DA), Approximate Karatsuba Multipliers (AKM), and Variable-Length Carry Select Adders (VLCSA). These results help validate that the architectural changes bring meaningful improvements in logic utilization and timing closure.

Furthermore, this chapter is structured to first present an overview of the results and then offer interpretation, followed by a multi-dimensional evaluation based on environmental, ethical, and economic factors. Simulation waveforms and internal RTL diagrams are also included to visually support the analysis.

5.2 Overview of Results

Post-synthesis and simulation results are collected and summarized to show the difference in performance between the two design iterations. The proposed design achieves better area efficiency and timing, but with a slight increase in dynamic power. This trade-off is analyzed carefully to understand the performance-cost balance.

The table below captures key design metrics, including cell usage, power, timing slack, and fanout. These parameters offer insights into both physical implementation characteristics and functional reliability. The use of approximate arithmetic modules played a significant role in achieving the optimizations

seen in the proposed implementation.

Overall, the results demonstrate a more balanced and scalable FIR architecture in the updated design. Its smaller area footprint and improved timing slack indicate readiness for ASIC/FPGA deployment in embedded systems where speed and area are key considerations.

5.3 Analysis and Interpretation

The proposed design improves significantly over the earlier version in terms of area and timing. The inclusion of DA reduces the dependency on real-time multipliers, thus minimizing logic gate usage. This leads to fewer combinational paths, which contributes to a reduction in propagation delay and improved slack margins.

Although the dynamic power increases slightly, it is largely due to enhanced parallel processing and the use of high-speed VLCSA structures. The increase is still within acceptable bounds for embedded DSP applications. In fact, the added power supports improved throughput and ensures results are generated every clock cycle once the pipeline is filled.

The design now supports better modularity and reuse. Functional blocks like AKM and VLCSA can be scaled or swapped with minimal changes. Such a structure is highly favorable in system-on-chip (SoC) environments where hardware reuse, area constraints, and timing compliance are critical factors.

5.4 Evaluation of Quality Factors

Environment: The proposed design reduces silicon area and leakage power, contributing positively toward energy-efficient computing. Smaller area and fewer interconnects mean lower electromagnetic interference and energy dissipation, which is beneficial in portable or eco-sensitive applications.

Sustainability: The modular and scalable nature of the design ensures it can be adapted to future projects or modified with minimal effort. Sustainability also applies in terms of lifecycle cost, as the reduced area and logic count can reduce long-term manufacturing and power costs.

Safety: Improved timing slack and lower fanout reduce the likelihood of glitches or metastability issues. These factors improve the safety and robustness of the circuit, especially in real-time control or medical systems where accuracy and reliability are paramount.

Ethics: The design process involved no use of restricted IPs or personal data. All synthesis was done using verifiable tools, and testing followed clean digital design practices. This ensures adherence to both academic and industry ethical guidelines.

Cost: Area reduction has a direct impact on silicon cost in both FPGA and ASIC contexts. With fewer resources consumed and optimized throughput, the design becomes a cost-effective solution. Higher dynamic power is offset by gains in speed and efficiency.

Type and Standards: The results are consistent with FIR filter performance expectations. The architecture adheres to design standards such as synchronous reset handling, clocked pipelining, and proper state machine control, ensuring compatibility with wider digital system design flows.

5.5 Waveform Simulation Results

Simulation waveforms are used to verify the behavior of the FIR filter with respect to clock and reset events, data input transitions, and final output values. The figures below show a clear correspondence between input patterns and the output after pipeline stages settle.




Figure 5.1: Simulation waveform of the earlier implementation

In the earlier implementation, the output shows delayed or undefined values during initial cycles due to incomplete pipeline control. The proposed design shows accurate and consistent output once the initial pipeline latency




Figure 5.2: Simulation waveform of the proposed implementation is overcome. This confirms the correct integration of the AKM, VLCSA, and DA stages in the architecture.

5.6 RTL Architecture Comparison

This section compares the RTL schematic structures of both implementations. RTL views help understand the datapath flow, signal routing, and module interconnections.

Earlier Design

The earlier design had a less organized layout and used more full-width logic for multiplication and accumulation. This made the design bulky and timing-critical, as reflected in the slack value of 0 ps.

Proposed Design

The proposed implementation shows a highly structured modular design. Each block (AKM, VLCSA, Register stages) is pipelined effectively. This enables continuous data flow, reduces critical paths, and ensures higher throughput. The design is also easier to debug and scale for more taps or higher bit-widths.

5.7 Conclusion

The proposed FIR filter design significantly improves over its predecessor in terms of area utilization, timing robustness, and modularity. Despite a modest increase in dynamic power, the gains in performance and synthesis




Figure 5.3: Approximate Karatsuba Multiplier – Earlier Design




Figure 5.4: VLCSA – Earlier Design




Figure 5.5: Complete Architecture – Earlier Design




Figure 5.6: Approximate Karatsuba Multiplier – Proposed Design




Figure 5.7: VLCSA – Proposed Design




Figure 5.8: Adder 4 Bit – Proposed Design




Figure 5.9: Complete Architecture – Proposed Design

efficiency make it a strong candidate for real-time signal processing in embedded platforms.

By using DA, AKM, and VLCSA, the design demonstrates that combining architectural and arithmetic optimizations can lead to real-world improvements in digital filter implementations. The results, verified through synthesis reports and waveform simulations, confirm both the functional correctness and performance alignment with design objectives.

In summary, the design strikes a well-balanced trade-off between speed, power, and cost—validating its effectiveness and viability for deployment in practical DSP applications such as audio filtering, wireless communication, and low-power wearable devices.

CHAPTER 6

Conclusions and Future Scope

6.1 Conclusions

This project presented the design and implementation of a high-performance, low-power FIR filter architecture using Distributed Arithmetic (DA), Approximate Karatsuba Multipliers (AKM), and Variable-Length Carry Select Adders (VLCSA). The primary goal was to optimize computational efficiency and resource utilization in real-time digital signal processing applications.

Through Verilog HDL design, simulation, and synthesis, it was demonstrated that the proposed architecture significantly improves timing slack, reduces logic complexity, and lowers area utilization compared to the conventional design. Although the dynamic power consumption increased slightly, this trade-off was acceptable in exchange for improved speed and scalability. Waveform validation and RTL analysis confirmed the correct functionality and structural improvements introduced by the modified datapath.

The integration of approximate computing (AKM) and optimized adder structures (VLCSA) with the DA framework proved to be an effective approach for enhancing FIR filter performance. The modularity and pipelined nature of the design make it suitable for scalable and portable embedded systems, validating the project's objectives and proving the success of the implemented optimizations.

6.2 Future Scope of Work

While the current implementation offers considerable improvements in speed and logic efficiency, there are still areas where enhancements can be made. Future work can explore incorporating dynamic coefficient update mechanisms to support adaptive filtering. This would make the architecture more versatile in applications like noise cancellation or real-time signal prediction.

In terms of hardware, applying advanced clock gating or power gating techniques could reduce the dynamic power further, especially in idle cycles. Additionally, migrating the design to lower technology nodes (e.g., 28nm or 14nm) or implementing it on low-power FPGAs may yield better power and performance figures.

On the algorithmic side, further exploration into alternative approximate multiplication techniques and hybrid adder designs could offer even more optimized trade-offs. The current design could also be extended to support higher filter orders or multi-channel processing with minimal architectural changes, making it suitable for more complex DSP tasks.

From a research perspective, integrating this FIR core into a larger DSP or SoC environment and evaluating its behavior under system-level constraints would be valuable. Real-time implementation and testing using tools like Xilinx Vivado or Intel Quartus on FPGA boards would help transition this architecture from theoretical to practical deployment.

REFERENCES

- [1] M. V. Pathi Amudalapalli, Nallagangula Mary Leena, S. V. Kiranmayi Sridhara, Lingam Sailaja, D. Bhavani, and Ch Kanakalingeswara Rao. “Implementation of FIR Filter Using Wallace Reduction Tree For High Speed Application”. In: *2024 First International Conference on Software, Systems and Information Technology (SSITCON)*. IEEE, 2024, pp. 1–6. DOI: 10.1109/SSITCON62437.2024.10795946.
- [2] N. Arumugam and B. Paramasivan. “An integrated FIR adaptive filter design by hybridizing canonical signed digit (CSD) and approximate booth recode (ABR) algorithm in DA architecture for the reduction of noise in the sensor nodes”. In: *Multidimensional Systems and Signal Processing* 32.4 (2021), pp. 1277–1311.
- [3] A. Agarwal and L. Bopanna. “Low latency area-efficient distributed arithmetic based multi-rate filter architecture for SDR receivers”. In: *Journal of Circuits, Systems and Computers* 27.08 (2018), p. 1850133.
- [4] S. Bose, A. De, and I. Chakrabarti. “Area-delay-power efficient VLSI architecture of FIR filter for processing seismic signal”. In: *IEEE Transactions on Circuits and Systems II: Express Briefs* 68.11 (2021), pp. 3451–3455.
- [5] Marco Castellano et al. “Algorithm and architecture for high speed merged arithmetic FIR filter generation”. In: *2008 3rd International Symposium on Communications, Control and Signal Processing*. IEEE. 2008.
- [6] Sandeep Kumar. “Design and Implementation of FIR Low Pass Filter Using Efficient Wallace Tree Multiplier”. In: *M.Tech Dissertation, National Institute of Technology Kurukshetra*. 2020.
- [7] K. Sakthi, J. Karthikeyan, and M. Jaya Surya. “Modified FIR Filter and High Speed Adder for Signal Processing Applications”. In: *International Conference on IoT Based Control Networks and Intelligent Systems (ICICNIS-2024)*. IEEE, 2024, pp. 581–583. DOI: 10.1109/ICICNIS64247.2024.10823281.
- [8] Sakkarai Samy Hari Krishnan and Krishnan Vidhya. “Distributed Arithmetic-FIR Filter Design Using Approximate Karatsuba Multiplier and VLCSA”. In: *Expert Systems with Applications* 249 (2024), p. 123488. DOI: 10.1016/j.eswa.2024.123488. URL: <https://doi.org/10.1016/j.eswa.2024.123488>.
- [9] Anton Blad and Oscar Gustafsson. “Redundancy reduction for high-speed FIR filter architectures based on carry-save adder trees”. In: *2010 IEEE International Symposium on Circuits and Systems*. IEEE. 2010.
- [10] V. Christilda and A. Milton. “Speed, power and area efficient 2D FIR digital filter using vedic multiplier with predictor and reusable logic”. In: *Analog Integrated Circuits and Signal Processing* 108.2 (2021), pp. 323–333.

- [11] M. M. A. da Rosa, G. Paim, P. U. L. da Costa, E. A. C. da Costa, R. I. Soares, and S. Bampi. “Axppa: Approximate parallel prefix adders”. In: *IEEE Transactions on Very Large Scale Integration (VLSI) Systems* 31.1 (2022), pp. 17–28.